

```
In [1]: from keras.datasets import cifar10
from keras import layers, models, optimizers
from keras.utils import to_categorical
from keras.preprocessing.image import ImageDataGenerator
```

```
In [2]: model = models.Sequential()
model.add(layers.Conv2D(32, (3, 3), activation= 'relu', input_shape= (32, 32, 3)))
model.add(layers.Conv2D(64, (3, 3), activation= 'relu'))
model.add(layers.MaxPooling2D((2,2)))
model.add(layers.Conv2D(128, (3,3), activation= 'relu'))
model.add(layers.Conv2D(128, (3,3), activation= 'relu'))
model.add(layers.MaxPooling2D((2,2)))
model.add(layers.Flatten())
model.add(layers.Dropout(0.5))
model.add(layers.Dense(512, activation= 'relu'))
model.add(layers.Dense(10, activation= 'softmax'))

model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
conv2d (Conv2D)	(None, 30, 30, 32)	896
conv2d_1 (Conv2D)	(None, 28, 28, 64)	18496
max_pooling2d (MaxPooling2D)	(None, 14, 14, 64)	0
conv2d_2 (Conv2D)	(None, 12, 12, 128)	73856
conv2d_3 (Conv2D)	(None, 10, 10, 128)	147584
max_pooling2d_1 (MaxPooling2D)	(None, 5, 5, 128)	0
flatten (Flatten)	(None, 3200)	0
dropout (Dropout)	(None, 3200)	0
dense (Dense)	(None, 512)	1638912
dense_1 (Dense)	(None, 10)	5130
=====		
Total params: 1,884,874		
Trainable params: 1,884,874		
Non-trainable params: 0		

```
In [3]: model.compile(optimizer= optimizers.RMSprop(lr=1e-4),
                    loss= 'categorical_crossentropy',
                    metrics= ['accuracy'])
```

```
In [4]: (x_train, y_train), (x_test, y_test) = cifar10.load_data()
```

```
In [5]: train_datagen = ImageDataGenerator(rescale= 1./255,
                                           rotation_range= 40,
                                           width_shift_range= 0.2,
                                           height_shift_range= 0.2,
                                           shear_range= 0.2,
                                           zoom_range= 0.2,
                                           horizontal_flip= True)
```

```
In [6]: x_test = x_test.astype('float32') / 255

y_train = to_categorical(y_train)
y_test = to_categorical(y_test)

x_val = x_train[:1000]
x_val = x_val.astype('float32')
x_part = x_train[1000:3000] # select small data set to see effect of data augment
y_val = y_train[:1000]
y_part = y_train[1000:3000]

train_generator = train_datagen.flow(x_part, y= y_part, batch_size= 20)
```


```
In [7]: history = model.fit(train_generator, steps_per_epoch= 100, epochs= 100, validation_steps= 100, validation_data=(x_val, y_val))

Epoch 83/100
100/100 [=====] - 3s 28ms/step - loss: 1.2942 - accuracy: 0.5270 - val_loss: 421.3049 - val_accuracy: 0.3390
Epoch 84/100
100/100 [=====] - 3s 28ms/step - loss: 1.2998 - accuracy: 0.5410 - val_loss: 500.6773 - val_accuracy: 0.3030
Epoch 85/100
100/100 [=====] - 3s 28ms/step - loss: 1.2917 - accuracy: 0.5415 - val_loss: 475.4423 - val_accuracy: 0.3270
Epoch 86/100
100/100 [=====] - 3s 28ms/step - loss: 1.3149 - accuracy: 0.5260 - val_loss: 462.2011 - val_accuracy: 0.3380
Epoch 87/100
100/100 [=====] - 3s 28ms/step - loss: 1.3038 - accuracy: 0.5475 - val_loss: 468.8743 - val_accuracy: 0.3350
Epoch 88/100
100/100 [=====] - 3s 28ms/step - loss: 1.3038 - accuracy: 0.5370 - val_loss: 327.4172 - val_accuracy: 0.3950
Epoch 89/100
100/100 [=====] - 3s 28ms/step - loss: 1.2764 - accuracy: 0.5475 - val_loss: 327.4172 - val_accuracy: 0.3950
```


```
In [8]: import matplotlib.pyplot as plt

acc = history.history['accuracy']
val_acc = history.history['val_accuracy']
loss = history.history['loss']
val_loss = history.history['val_loss']
epochs = range(1, len(acc) + 1)

plt.plot(epochs, loss, 'bo', label= 'Training Loss')
plt.plot(epochs, val_loss, 'r', label= 'Validation Loss')
plt.title('Training and validation loss of cifar10 model')
plt.legend()
plt.savefig('Results/6_2b/Loss.png');
```



```
In [10]: plt.plot(epochs, acc, 'bo', label= 'Training accuracy')
plt.plot(epochs, val_acc, 'r', label= 'Validation accuracy')
plt.title('Training and validation accuracy of cifar10 model')
plt.legend()
plt.savefig('Results/6_2b/Accuracy.png');
```



```
In [11]: model.save('Results/6_2b/model2.h5')
```

```
In [12]: result = model.evaluate(x_test, y_test)
print(f'The model result for the test data is loss {result[0]} and accuracy {result[1]}')

313/313 [=====] - 3s 11ms/step - loss: 1.3944 - accuracy: 0.5337
The model result for the test data is loss 1.3944133520126343 and accuracy 0.5336999893188477
```

```
In [13]: import numpy as np
import pandas as pd
preds = model.predict(x_test)
pred_class = np.argmax(preds, axis= 1)
act = np.argmax(y_test, axis= 1)
pred_df = pd.DataFrame({'Actual':act, 'Predicted':pred_class})
pred_df
```

Out[13]:

	Actual	Predicted
0	3	3
1	8	1
2	8	8
3	0	0
4	6	4
...
9995	8	5
9996	3	6
9997	5	5
9998	1	1
9999	7	7

10000 rows × 2 columns

```
In [14]: pred_df.to_csv('Results/6_2b/predictions.csv')
```

In []: